



Class
Recommendation Engines



Topic
Introduction

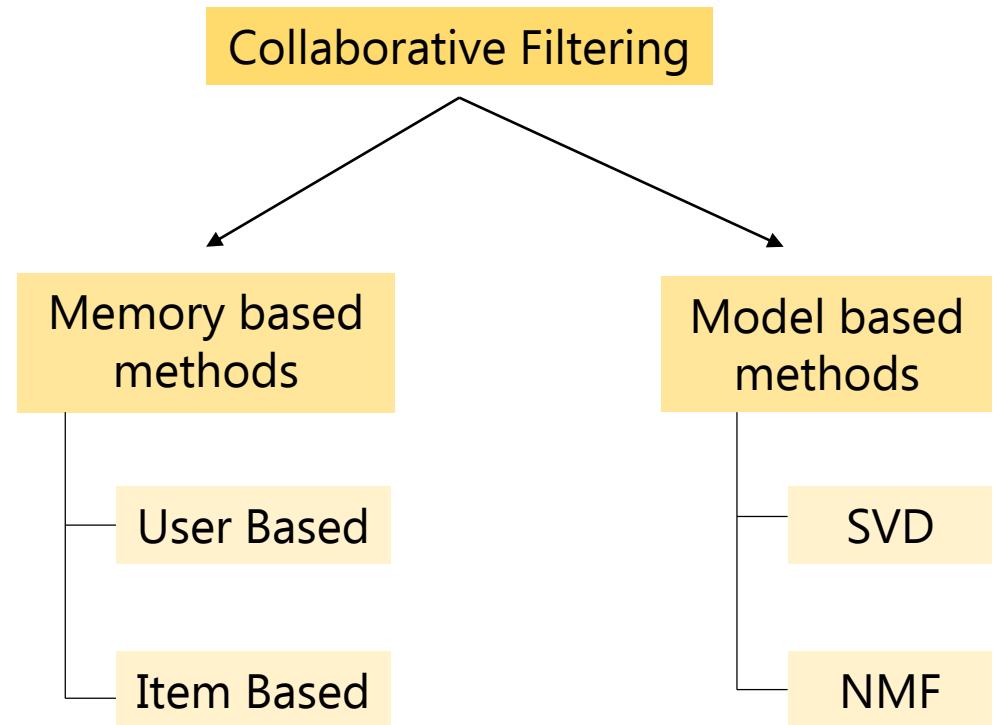
Collaborative Filtering

There are many ways to create recommendation engines

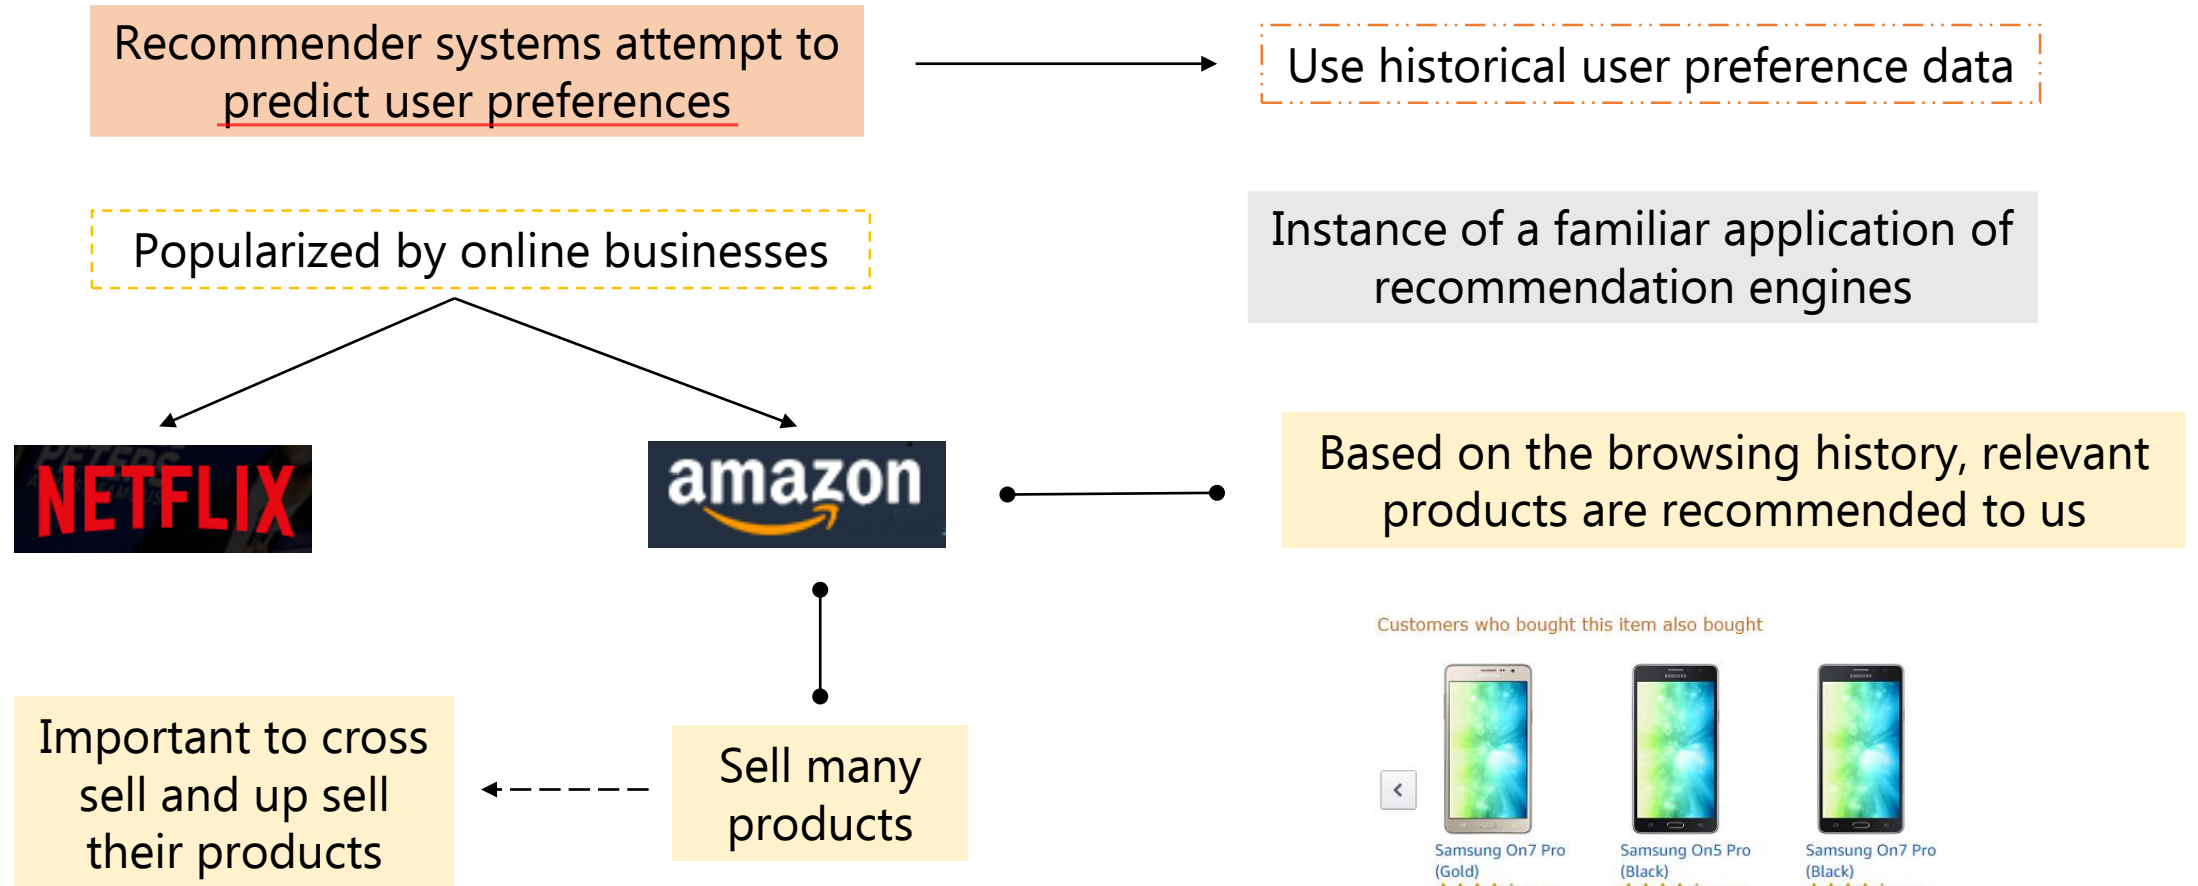


Classified as an unsupervised or semi-supervised learning algorithm

Collaborative filtering is necessarily unsupervised



Predicting User Preferences



Collaborative Filtering

User item rating matrix

	Item 1	Item 2	Item 3	Item 4	Item 5	Item 6	Item 7
User 1	3	5	1	2	4	3	3
User 2	2	4	1	2	?	3	2
User 3	3	?	5	?	4	1	1
User 4	4	5	1	?	?	?	?

Ratings for a user item pair



Items - movies or TV series

User item rating matrix - user movie rating matrix

Collaborative filtering works on the premise

Users who are similar, select similar kind of products

Gauging how similar a product is to a set of products, relevant recommendations are made



Collaborative Filtering

User item rating matrix

	Item 1	Item 2	Item 3	Item 4	Item 5	Item 6	Item 7
User 1	3	5	1	2	4	3	3
User 2	2	4	1	2	?	3	2
User 3	3	?	5	?	4	1	1
User 4	4	5	1	?	?	?	?

Predicted value of these ratings of an item is high enough



The item is recommended to the user



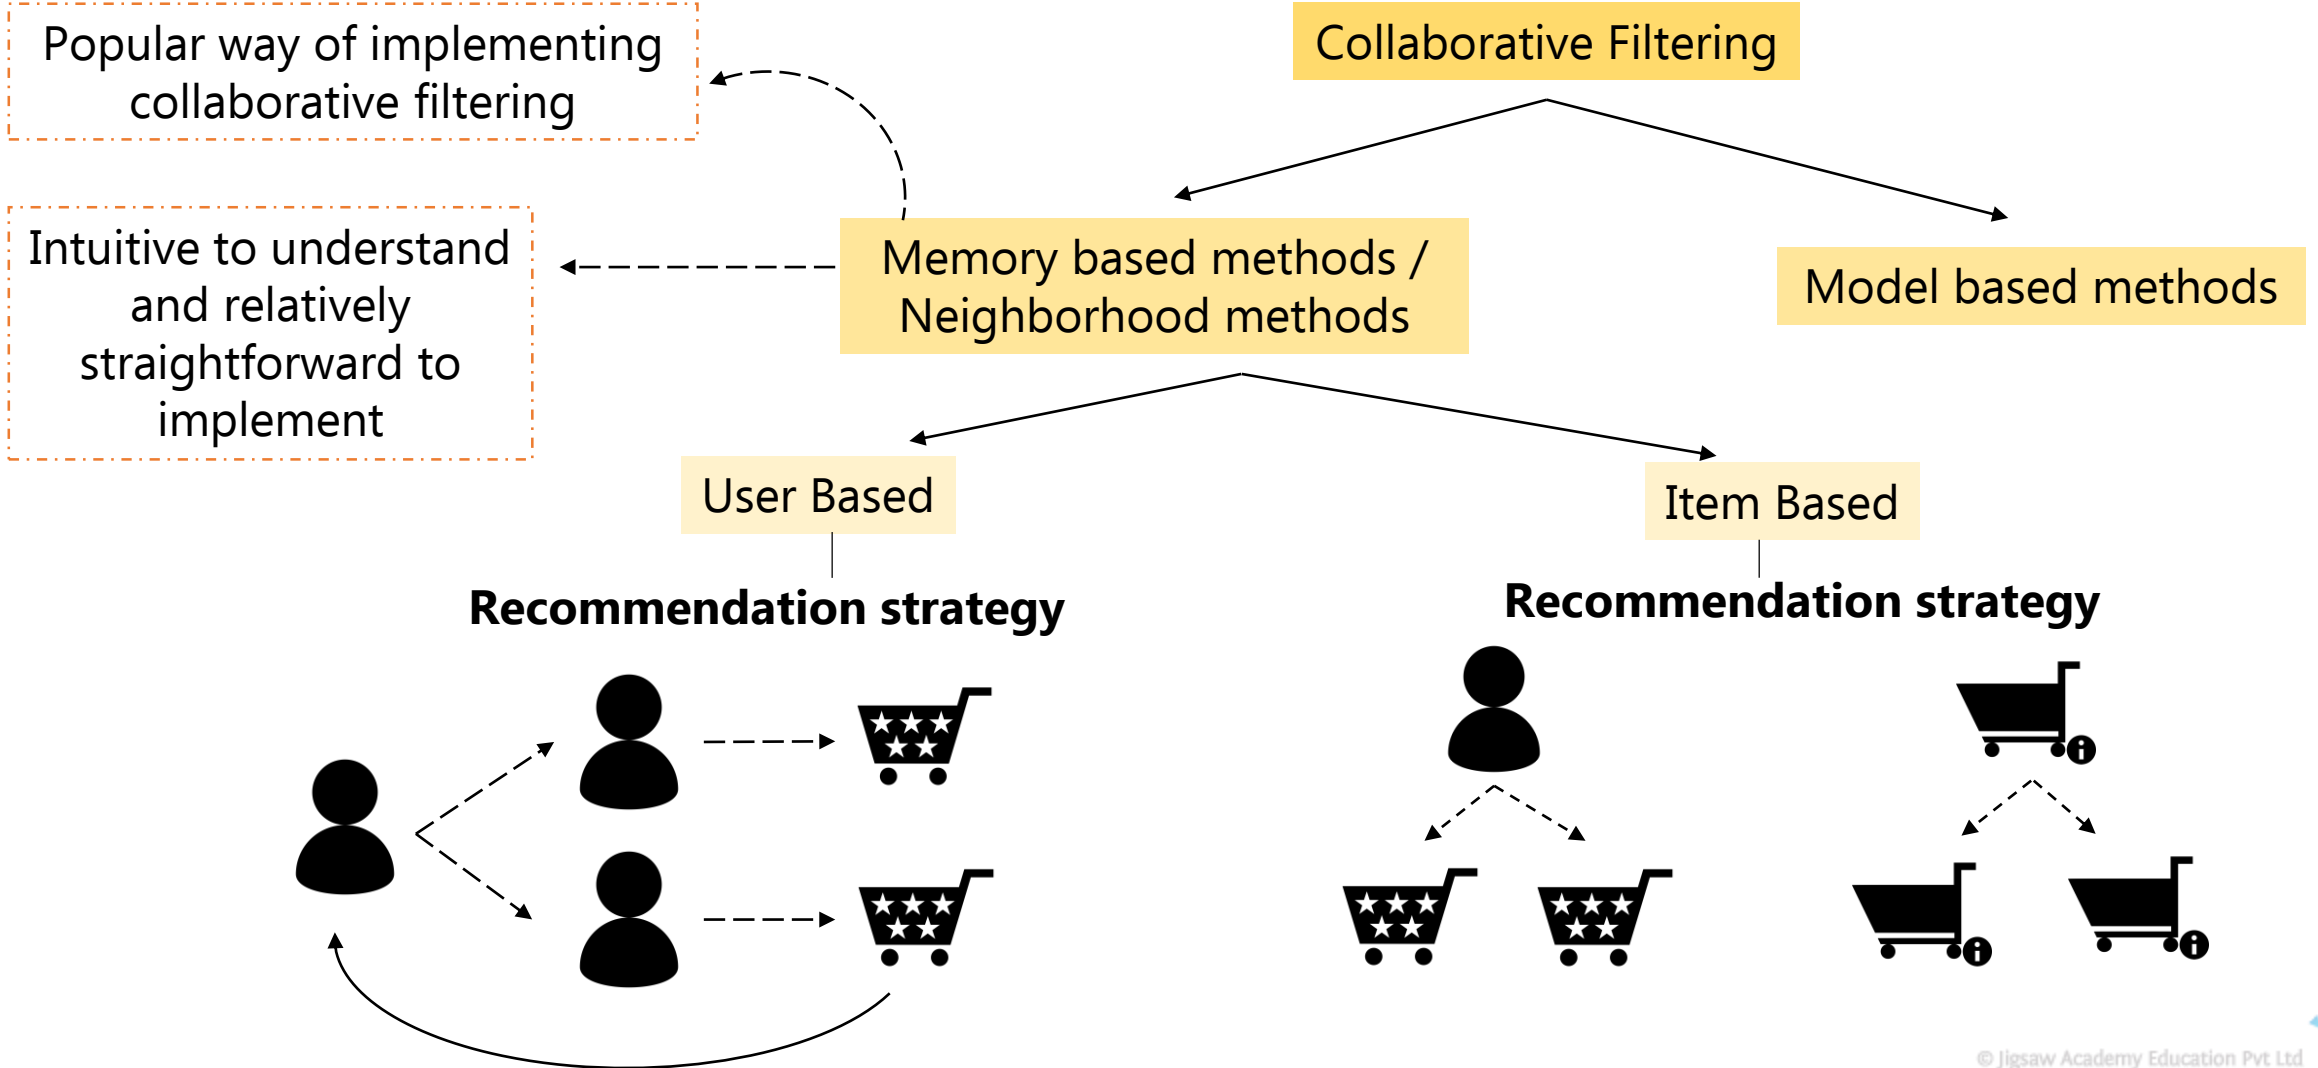
Usually a **user item rating matrix** is **partially populated** as not all users tend to rate all items



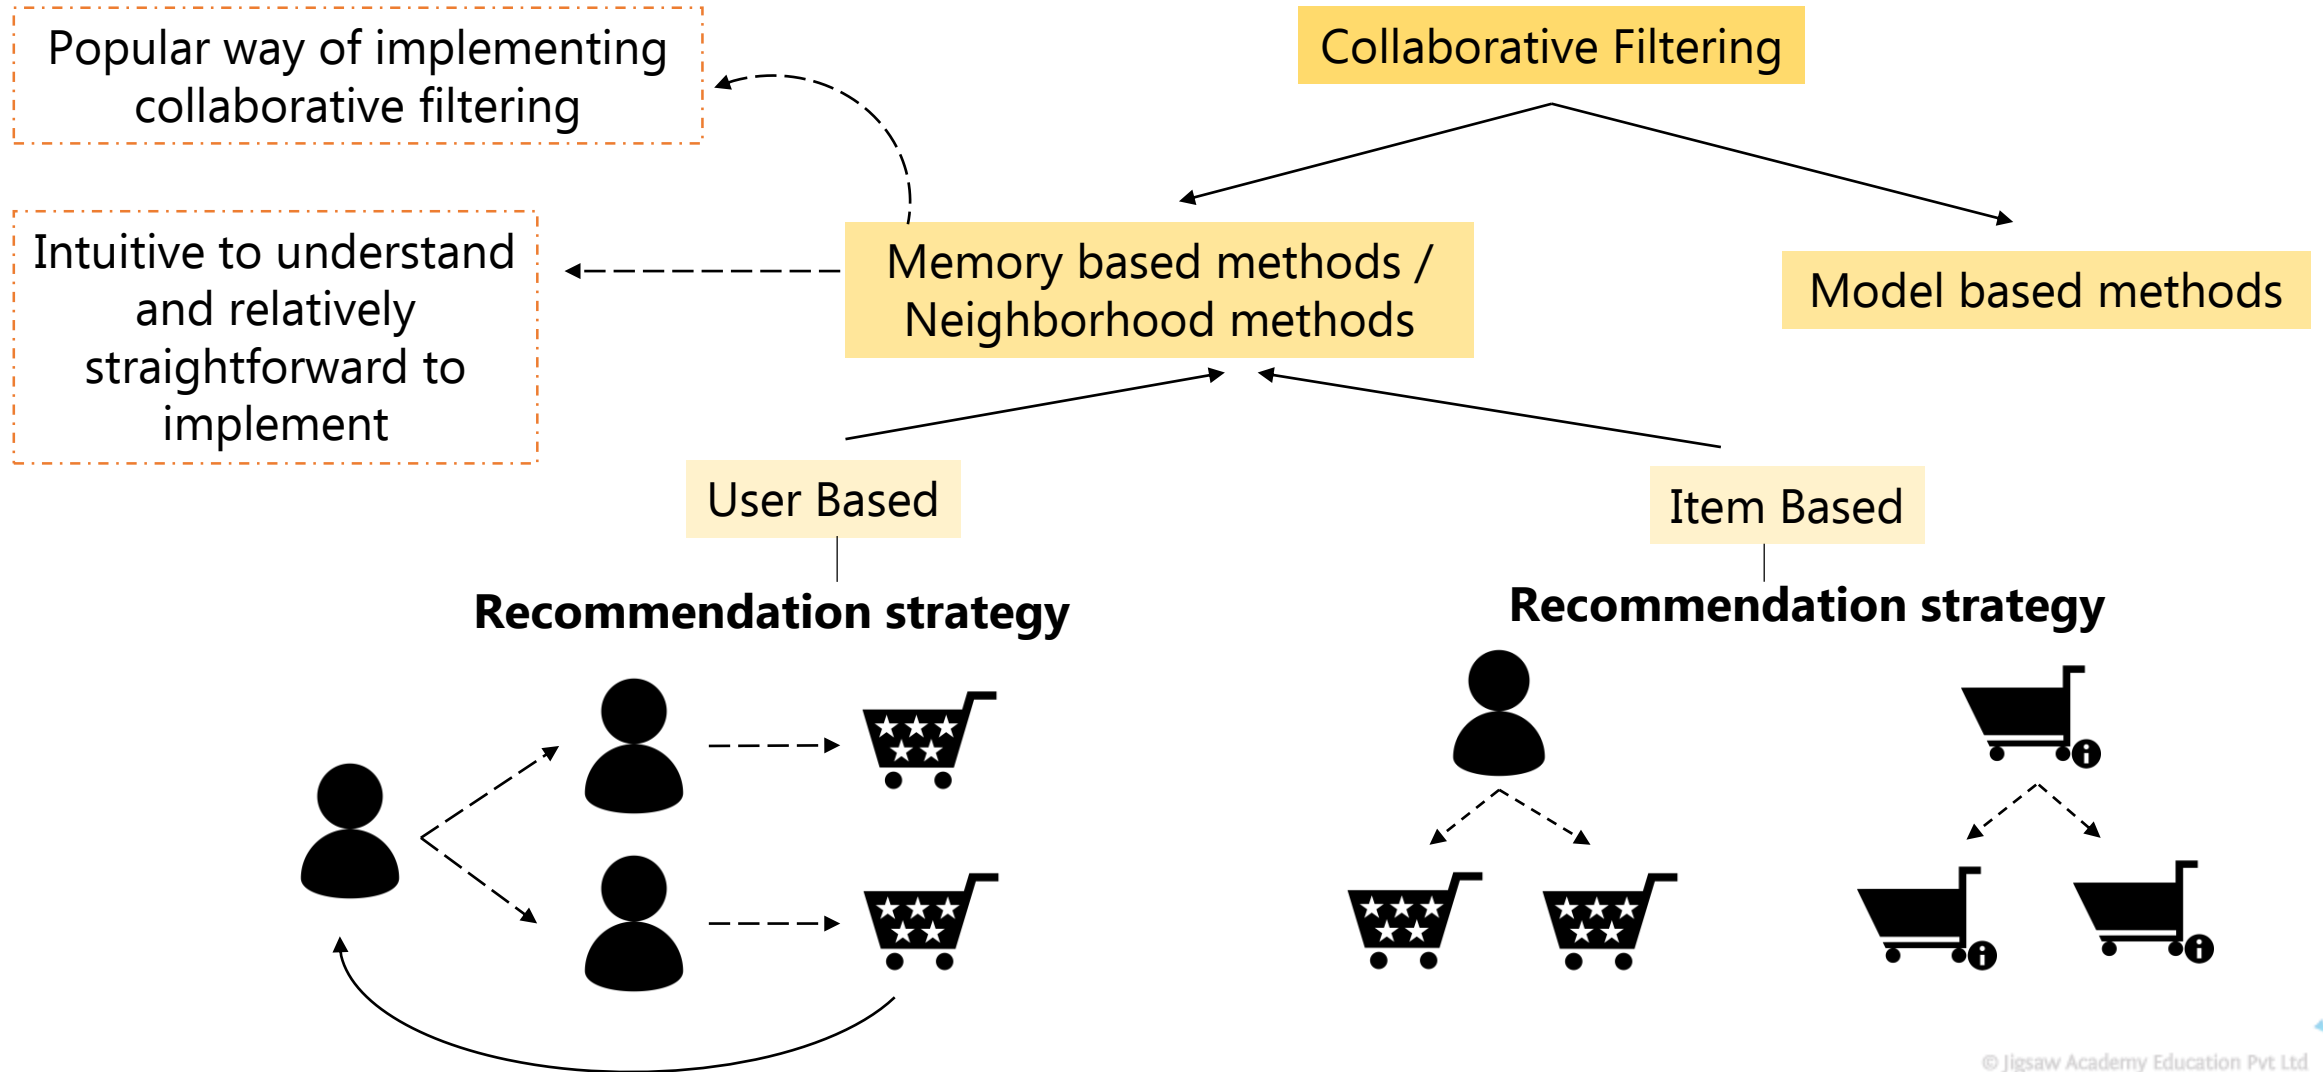
Recommendation engines usually try to predict these missing ratings



Memory Based/ Neighbourhood Methods



Memory Based/ Neighbourhood Methods



User Based Collaborative Filtering

User item rating matrix

	Item 1	Item 2	Item 3	Item 4	Item 5
Alice	5	3	4	4	?
User 1	3	1	2	3	3
User 2	4	3	4	3	5
User 3	3	3	1	5	4
User 4	1	5	5	2	1

What is the rating for the Item 5 given by Alice?

Find users similar to Alice

Use any measures of similarity such as a **Pearson co-relation** or **Cosine similarity**

Based on the most similar users

Predict the rating for item 5 by Alice



User Based Collaborative Filtering

User item rating matrix

	Item 1	Item 2	Item 3	Item 4	Item 5
Alice	5	3	4	4	?
User 1	3	1	2	3	3
User 2	4	3	4	3	5
User 3	3	3	1	5	4
User 4	1	5	5	2	1

The rating for item 5 by Alice will predicted based on how User 1 and User 2 have rated item 5

What is the rating for the Item 5 given by Alice?

Find users similar to Alice

Use any measures of similarity such as a **Pearson co-relation** or **Cosine similarity**

User 1 and User 2 are most similar users to Alice



Recap

- Predicting user preferences
- Collaborative filtering
- Memory based methods
- User based collaborative filtering





Class

Recommendation Engines

★ ★

Topic

Item Based Collaborative Filtering

Item Based Collaborative Filtering

User item rating matrix

	Item 1	Item 2	Item 3	Item 4	Item 5
Alice	5	3	4	4	?
User 1	3	1	2	3	3
User 2	4	3	4	3	5
User 3	3	3	1	5	4
User 4	1	5	5	2	1

What is the rating for the Item 5 given by Alice?

Find out the items similar to item 5 that are rated by Alice

Use either Co-relation or Cosine similarity to find the similar items

Based on the items that are more similar to item 5, rating of item 5 by Alice will be computed



Item Based Collaborative Filtering

User item rating matrix

	Item 1	Item 2	Item 3	Item 4	Item 5
Alice	5	3	4	4	?
User 1	3	1	2	3	3
User 2	4	3	4	3	5
User 3	3	3	1	5	4
User 4	1	5	5	2	1

What is the rating for the Item 5 given by Alice?

Find out the items similar to item 5 that are rated by Alice

Item 1 and Item 4 turn out to be the most similar items to Item 5

Based on Alice's rating of Item 1 and Item 4, the prediction about the rating for Item 5 would be made



Code Demo



Prediction



Prediction in a Recommender system, particularly the memory based Recommender systems, can be done in 2 ways:

1. Predict the rating based on just the ratings and similarities of items or users

$$\hat{r}_{ui} = \frac{\sum_{v \in N_i^k(u)} \text{sim}(u, v) \cdot r_{vi}}{\sum_{v \in N_i^k(u)} \text{sim}(u, v)}$$

$$\hat{r}_{ui} = \frac{\sum_{j \in N_u^k(i)} \text{sim}(i, j) \cdot r_{uj}}{\sum_{j \in N_u^k(i)} \text{sim}(i, j)}$$

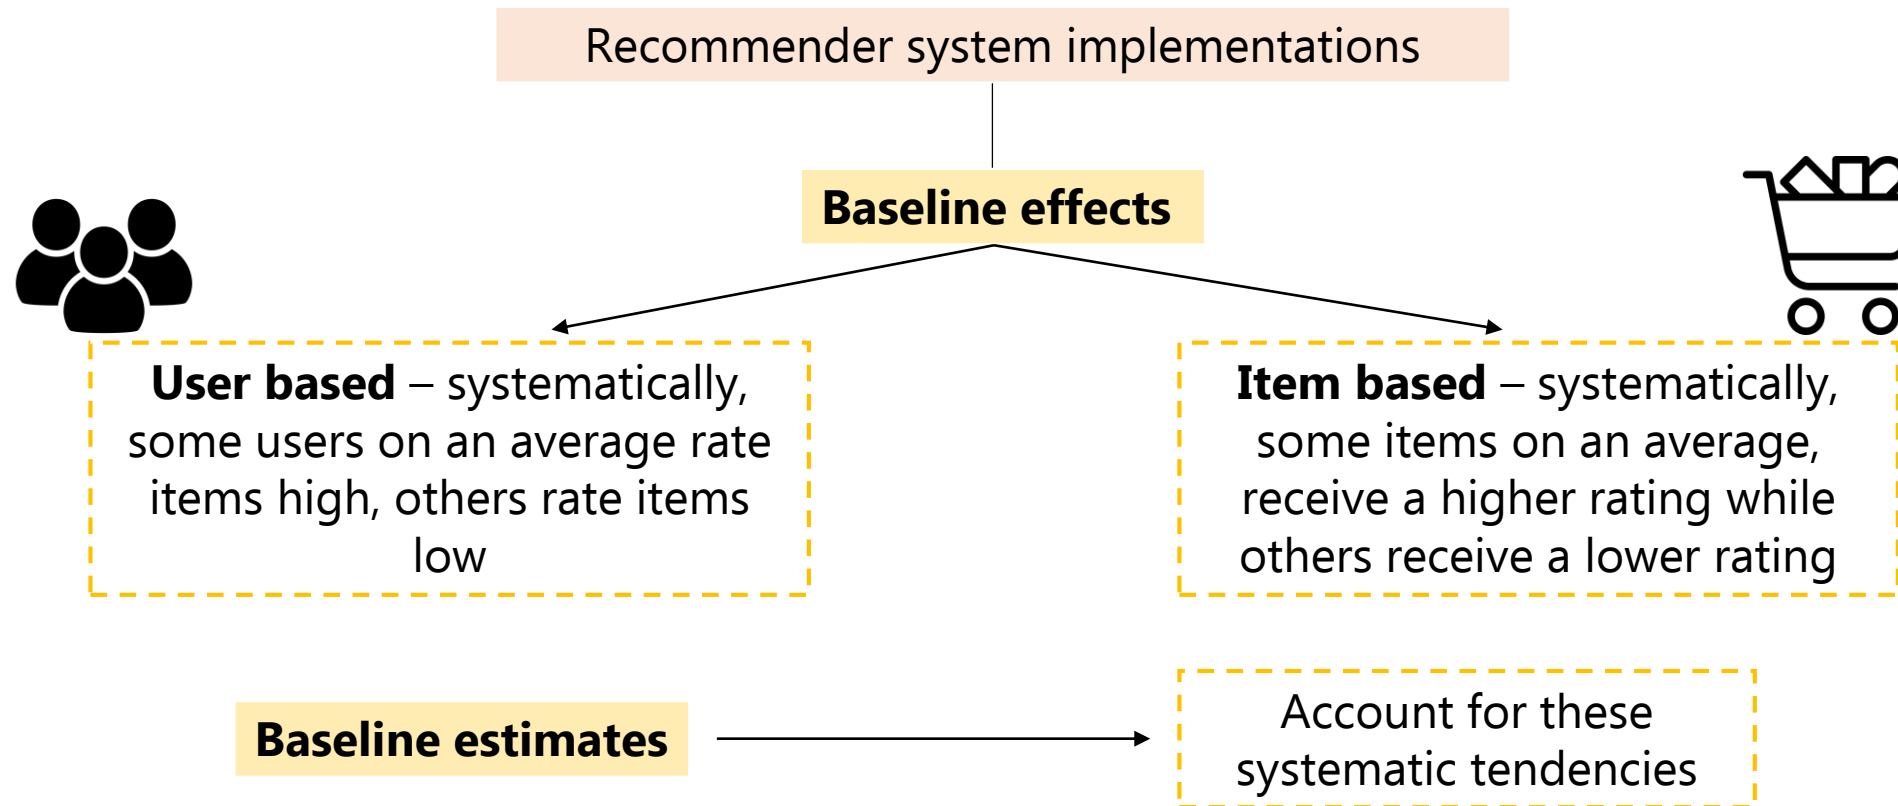
2. Predict the rating based on the average effects of ratings of items and users as well

$$\hat{r}_{ui} = \mu_i + \frac{\sum_{j \in N_u^k(i)} \text{sim}(i, j) \cdot (r_{uj} - \mu_j)}{\sum_{j \in N_u^k(i)} \text{sim}(i, j)}$$

$$\hat{r}_{ui} = \mu_u + \frac{\sum_{v \in N_i^k(u)} \text{sim}(u, v) \cdot (r_{vi} - \mu_v)}{\sum_{v \in N_i^k(u)} \text{sim}(u, v)}$$




Baseline Effects



Not all software libraries have the ability to take into account the baseline effects

Accuracy




The accuracy of a Recommender system can be measured by estimating the metrics




Root Mean Squared Error (RMSE)

Mean Absolute Error (MAE)



Computation of these quantities is quite straightforward since the actual ratings for many user and item pairs are already known

Hyperparameters




What would be the hyperparameters for neighborhood based recommendation systems?

1. Number of neighbors

2. Similarity metric

3. Prediction method

Grid Search using K-fold Cross Validation to figure out which values of these hyperparameters will be most suitable



Recap

- Item based collaborative filtering
- Code demo of item based collaborative filtering
- Prediction
- Baseline effects
- Accuracy
- Hyperparameters



Class
Recommendation Engines



Topic
Model Based

Matrix Factorization




Build recommendation engines using - user based and item based collaborative filtering

Another class of collaborative filtering recommendation engines that rely on the use



Matrix factorization algorithms



Singular Value Decomposition

Non Negative Matrix Factorization

Matrix Factorization

User item rating matrix

	Star Trek	Avatar	Spiderman	Hulk
User 1	4	4	1	1
User 2	5	5	2	2
User 3	1	1	4	4
User 4	2	2	5	5

Assume –
Ratings are on a scale of 1 to 5



In context of recommendation engines, matrix factorization indicates finding 2 matrices :

User factor matrix



Item factor matrix



Matrix Factorization

User item rating matrix

	Star Trek	Avatar	Spiderman	Hulk
User 1	4	4	1	1
User 2	5	5	2	2
User 3	1	1	4	4
User 4	2	2	5	5

Factorizing

User factor matrix

Dim 1	Dim 2
0.5	-0.41
0.5	-0.57
-0.5	0.41
-0.5	0.57

Item factor matrix

	Star Trek	Avatar	Spiderman	Hulk
Dim 1	0.5	0.5	-0.5	-0.5
Dim 2	-0.2	-0.2	0.7	0.7



What does each
of these
matrices signify?



Matrix Factorization

User item rating matrix

	Star Trek	Avatar	Spiderman	Hulk
User 1	4	4	1	1
User 2	5	5	2	2
User 3	1	1	4	4
User 4	2	2	5	5

Genre: **Science Fiction (Sci-Fi)**

Genre: **Fantasy movies**

User factor matrix

Dim 1	Dim 2
0.5	-0.41
0.5	-0.57
-0.5	0.41
-0.5	0.57

2 columns

Representing factors

Movie genres

Item factor matrix

	Star Trek	Avatar	Spiderman	Hulk
Dim 1	0.5	0.5	-0.5	-0.5
Dim 2	-0.2	-0.2	0.7	0.7



Matrix Factorization

User item rating matrix

	Star Trek	Avatar	Spiderman	Hulk
User 1	4	4	1	1
User 2	5	5	2	2
User 3	1	1	4	4
User 4	2	2	5	5

High values
corresponding
to Sci-Fi genre

User factor matrix

Sci-Fi	Dim 2
0.5	-0.41
0.5	-0.57
-0.5	0.41
-0.5	0.57

High rating to Sci-Fi
movies in the
original matrix

Numbers represent
how much each factor,
each of the items have

Item factor matrix

	Star Trek	Avatar	Spiderman	Hulk
Dim 1	0.5	0.5	-0.5	-0.5
Dim 2	-0.2	-0.2	0.7	0.7



Matrix Factorization

User item rating matrix

	Star Trek	Avatar	Spiderman	Hulk
User 1	4	4	1	1
User 2	5	5	2	2
User 3	1	1	4	4
User 4	2	2	5	5

User factor matrix

Sci-Fi	Dim 2
0.5	-0.41
0.5	-0.57
-0.5	0.41
-0.5	0.57

Item factor matrix

	Star Trek	Avatar	Spiderman	Hulk
Sci-Fi	0.5	0.5	-0.5	-0.5
Dim 2	-0.2	-0.2	0.7	0.7



Matrix Factorization

User item rating matrix

	Star Trek	Avatar	Spiderman	Hulk
User 1	4	4	1	1
User 2	5	5	2	2
User 3	1	1	4	4
User 4	2	2	5	5

User factor matrix

Sci-Fi	Fantasy
0.5	-0.41
0.5	-0.57
-0.5	0.41
-0.5	0.57

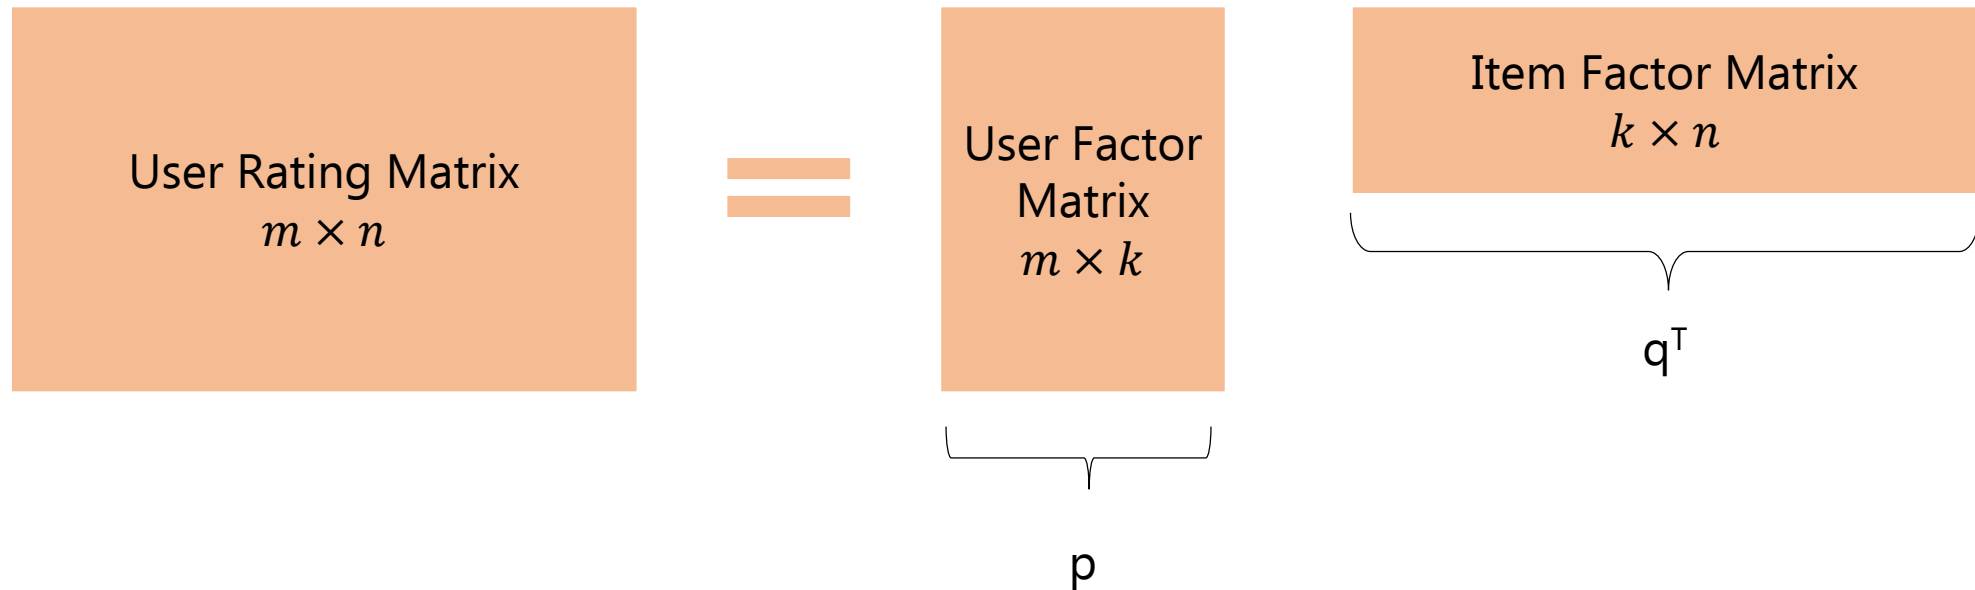
Item factor matrix

	Star Trek	Avatar	Spiderman	Hulk
Sci-Fi	0.5	0.5	-0.5	-0.5
Fantasy	-0.2	-0.2	0.7	0.7



Matrix Factorization

Technical details of matrix factorization

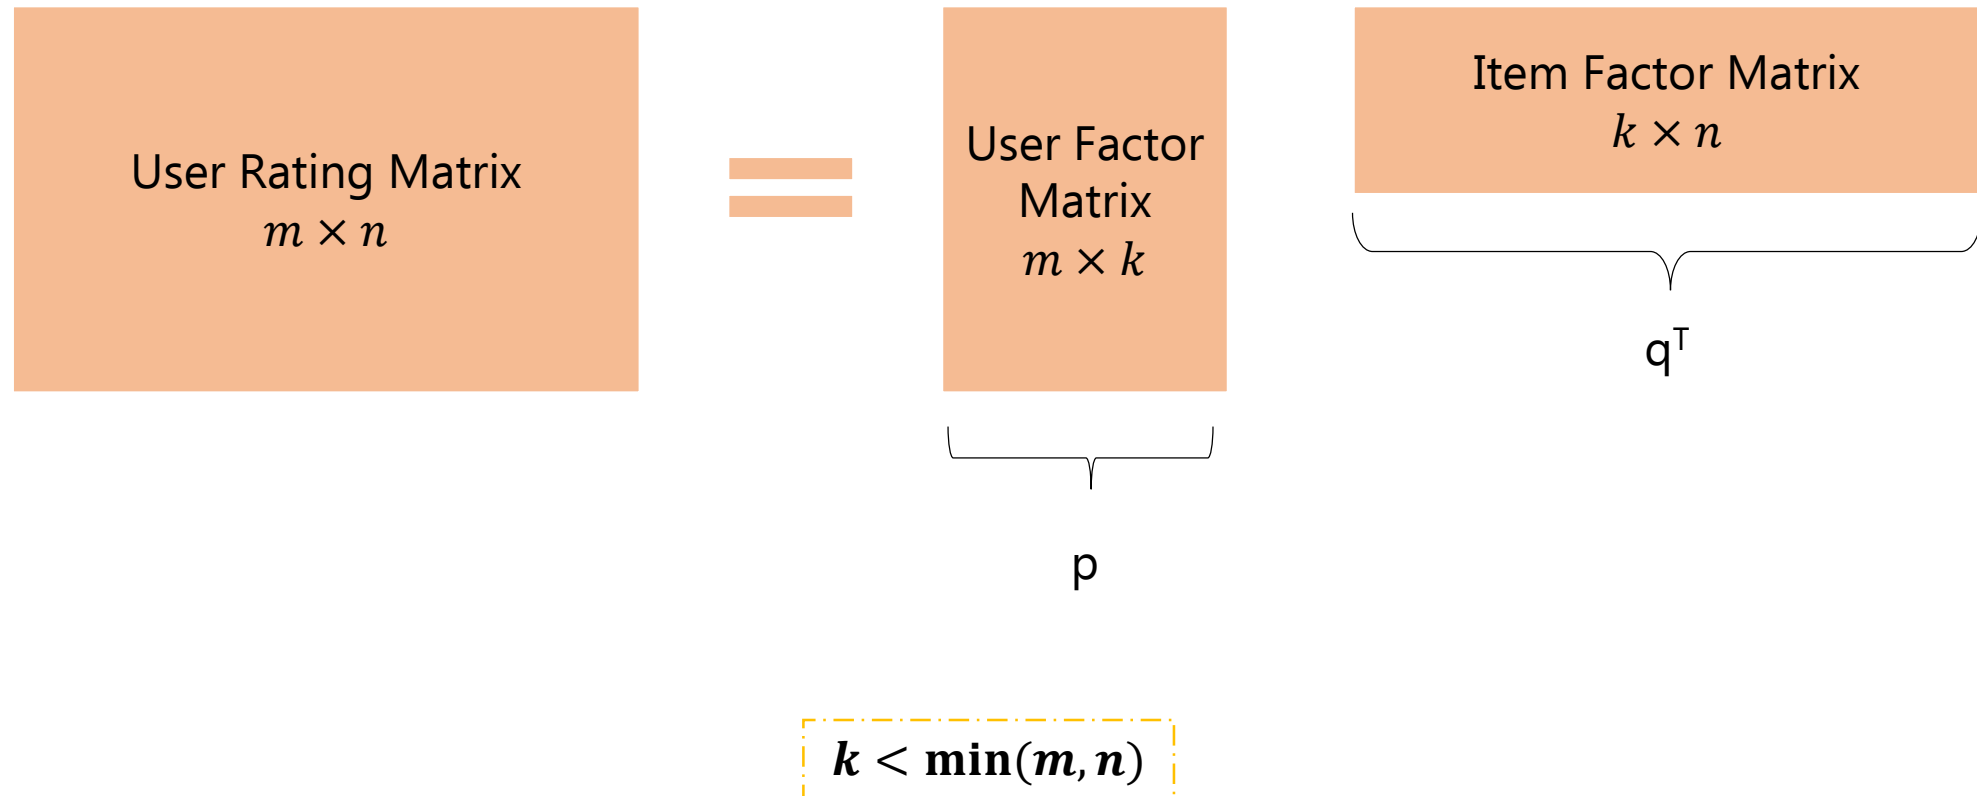


k is determined by the user

Number of factors to be considered is always less than the number of rows, or columns or original rating matrix, whichever is minimum

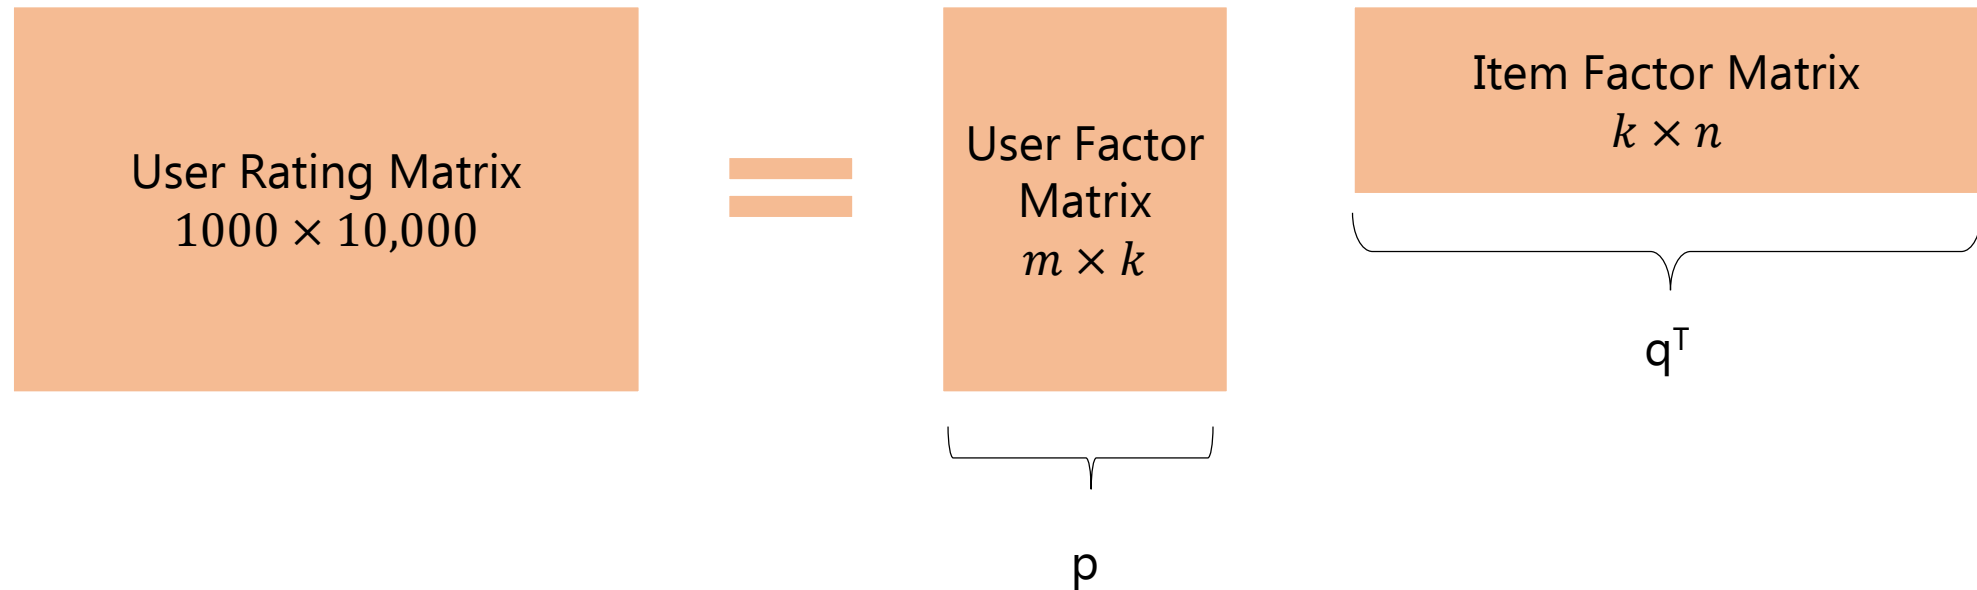
Matrix Factorization

Technical details of matrix factorization



Matrix Factorization

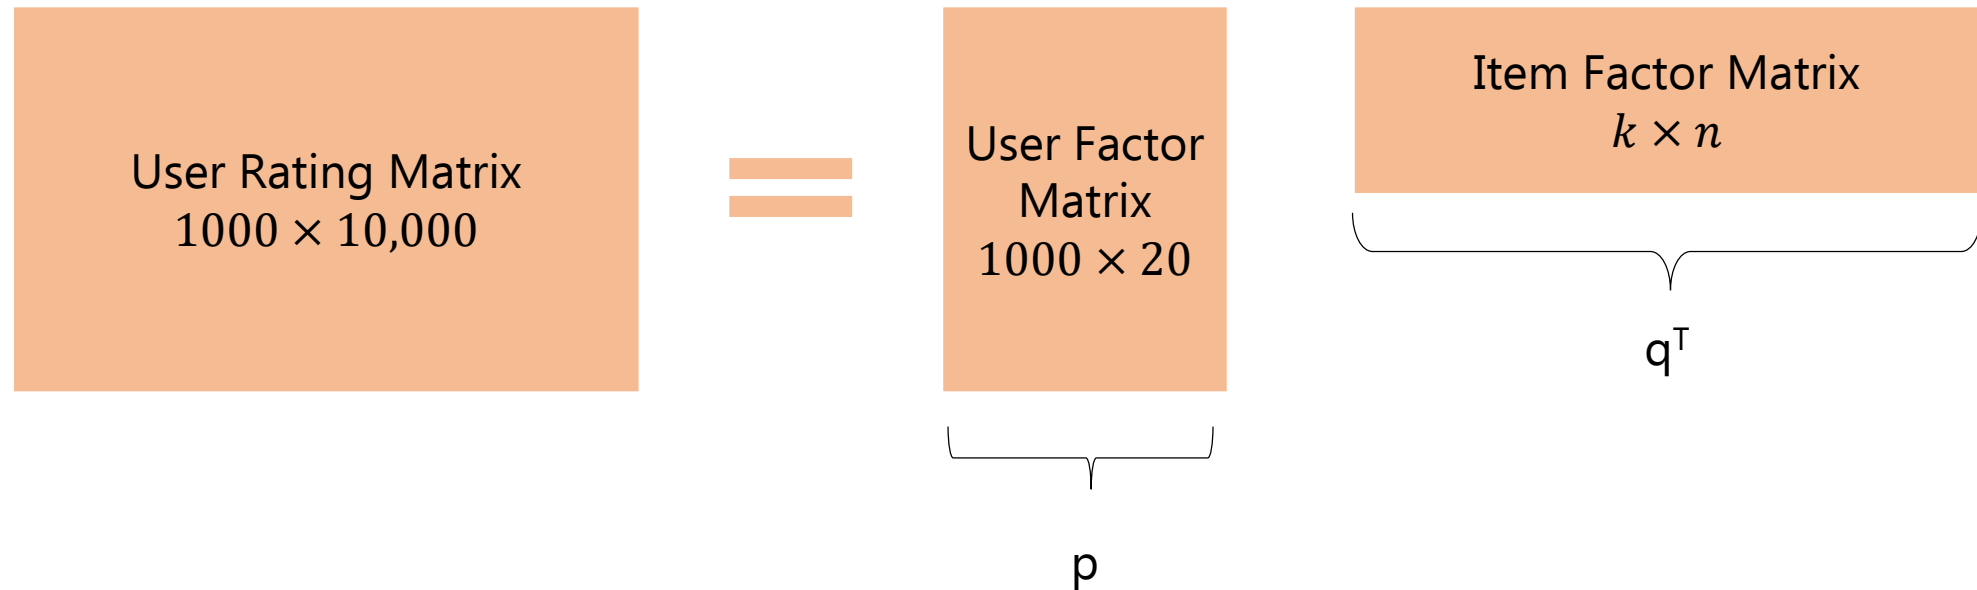
Technical details of matrix factorization



$$20 < \min(1000, 10000)$$

Matrix Factorization

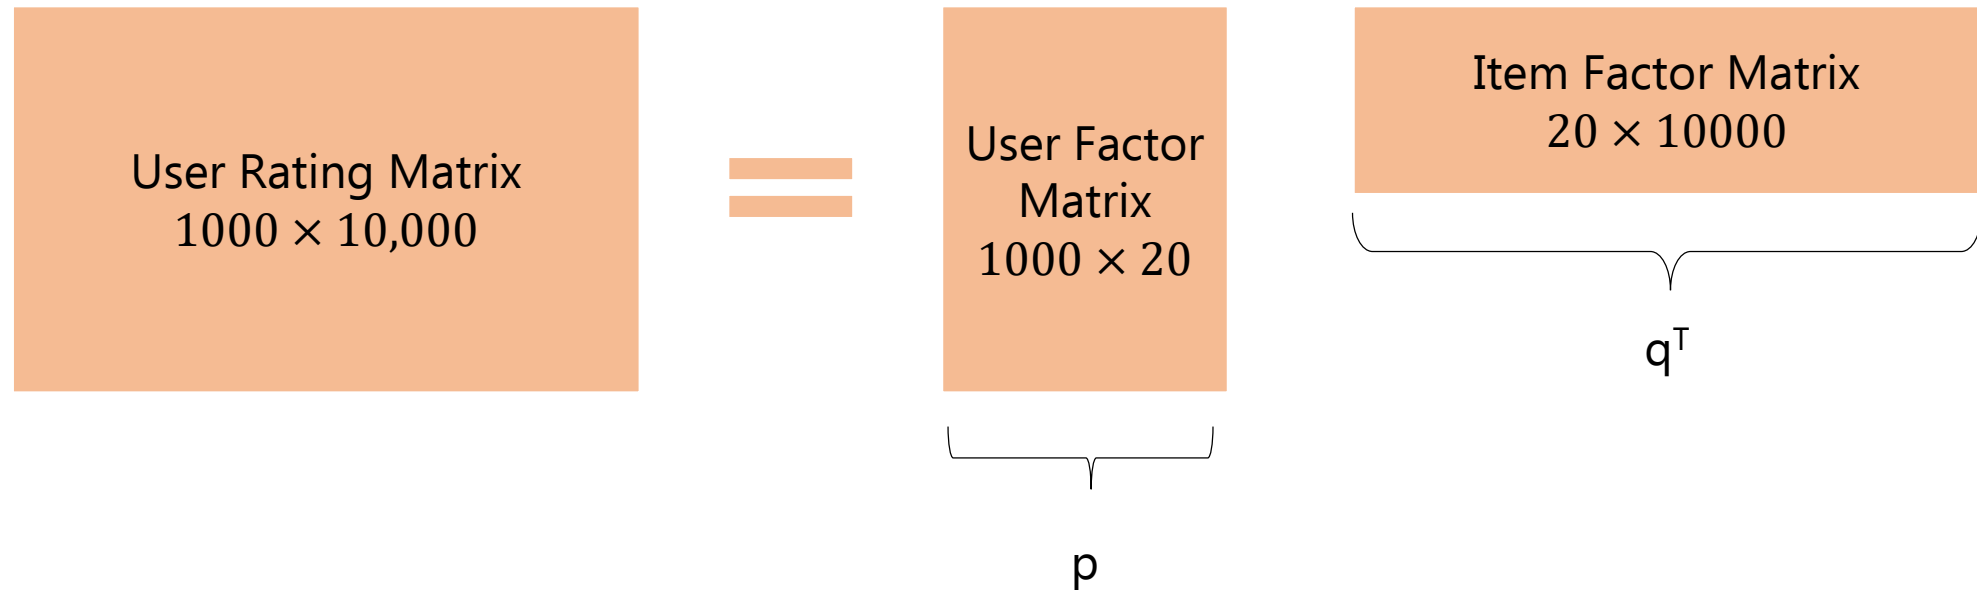
Technical details of matrix factorization



$$20 < \min(1000, 10000)$$

Matrix Factorization

Technical details of matrix factorization



$$20 < \min(1000, 10000)$$

Recap

- Matrix factorization

Class
Recommendation Engines



Topic
**Model Based Methods: Prediction and Estimation
Using SVD and NMF**

Matrix Factorization



$$\hat{r}_{ui} = \mu + b_{\mu} + b_i + p_u q_i^T$$

μ – Average of all the ratings in the data

b_{μ} – Estimate of average ratings considering each user

b_i – Estimate of average ratings considering each item

p_u – User factor matrix obtained from matrix factorization

q_i^T – Item factor matrix



Matrix Factorization

$$\hat{r}_{ui} = \mu + b_{\mu} + b_i + p_u q_i^T$$

	Star Trek	Avatar	Spiderman	Hulk
User 1	4	?	3	?
User 2	4	?	2	1
User 3	3	?	4	3
User 4	?	?	?	?

Dim 1	Dim 2
0.8	0.4
0.6	-0.5
0.2	0.1

3×2

Each entry in this matrix talks about user preference for each of the 2 factors

What will be the form of User factor and Item factor matrix?

If a given item or user doesn't have any rating history, then that item or user is dropped before doing matrix factorization



Matrix Factorization

$$\hat{r}_{ui} = \mu + b_{\mu} + b_i + p_u q_i^T$$

	Star Trek	Avatar	Spiderman	Hulk	b_u
User 1	4	?	3	?	1.75
User 2	4	?	2	1	1.75
User 3	3	?	4	3	2.5
User 4	?	?	?	?	0
b_i	2.75	0	2.75	1	

b_i - Computed the column averages

b_u - Computed the row averages

Dim 1	Dim 2
0.8	0.4
0.6	-0.5
0.2	0.1

3×2

	Star Trek	Spiderman	Hulk
Dim 1	0.6	0.8	-0.4
Dim 2	0.2	0.3	0.1

2×3



Matrix Factorization

$$\hat{r}_{ui} = \mu + b_{\mu} + b_i + p_u q_i^T$$

	Star Trek	Avatar	Spiderman	Hulk	b_u
User 1	4	?	3	?	1.75
User 2	4	?	2	1	1.75
User 3	3	?	4	3	2.5
User 4	?	?	?	?	0
b_i	2.75	0	2.75	1	

p

Dim 1	Dim 2
0.8	0.4
0.6	-0.5
0.2	0.1

3×2

q^T

	Star Trek	Spiderman	Hulk
Dim 1	0.6	0.8	-0.4
Dim 2	0.2	0.3	0.1

2×3

Global average of ratings

$$\mu = 1.5$$

$$\hat{r}_{14} = 1.5 + 1.75 + 1 + p_1 q_4^T$$

$$p_1 q_4^T = [0.8, 0.4] \times [-0.4, 0.1]$$

$$p_1 q_4^T = [-0.32 + 0.8 - 0.16 + 0.4]$$

$$p_1 q_4^T = -0.36$$

$$\hat{r}_{14} = 1.5 + 1.75 + 1 - 0.36$$

$$\hat{r}_{14} = 3.89$$



Matrix Factorization

$$\hat{r}_{ui} = \mu + b_{\mu} + b_i + p_u q_i^T$$

$$\mu = 1.5$$

	Star Trek	Avatar	Spiderman	Hulk	b_u
User 1	4	?	3	?	1.75
User 2	4	?	2	1	1.75
User 3	3	?	4	3	2.5
User 4	?	?	?	?	0
b_i	2.75	0	2.75	1	

$$\hat{r}_{22} = 1.5 + 1.75 + 0 + p_2 q_2^T$$

$$p_2 q_2^T = [0.6, -0.5] \times [0, 0]$$

$$p_2 q_2^T = 0$$

$$\hat{r}_{22} = 1.5 + 1.75 + 0 + 0$$

$$\hat{r}_{22} = 3.25$$

Dim 1	Dim 2	p
0.8	0.4	
0.6	-0.5	
0.2	0.1	

3×2

	Star Trek	Spiderman	Hulk	q^T
Dim 1	0.6	0.8	-0.4	
Dim 2	0.2	0.3	0.1	

2×3



Matrix factorization



How is the matrix factorization done?

How does Singular Value Decomposition work?

It will be discussed in the context of recommendation engines with an emphasis on providing an intuition on how things work

2 popular algorithms

Singular Value Decomposition (SVD)

Non Negative Matrix Factorization (NMF)



Singular Value Decomposition

Rating matrix

	Star Trek	Avatar	Spiderman	Hulk
User 1	4	?	3	?
User 2	4	?	2	1
User 3	3	?	4	3
User 4	?	?	?	?



Singular Value Decomposition

Rating matrix

	Star Trek	Spiderman	Hulk
User 1	4	3	?
User 2	4	2	1
User 3	3	4	3

It will be imputed with either the mean or 0 value

Depending on which implementation is being used to do the matrix factorization

Some implementations may not even impute the values



Singular Value Decomposition


Rating matrix

	Star Trek	Spiderman	Hulk
User 1	4	3	3.5
User 2	4	2	1
User 3	3	4	3

$$\hat{r}_{ui} = \mu + b_{\mu} + b_i + p_u q_i^T$$

Predictions

	Star Trek	Spiderman	Hulk
User 1	3.2	2.9	3.3
User 2	3.8	2.5	1.2
User 3	2.9	3.8	3.4



$$Error = (4 - 3.2)^2 + (3 - 2.9)^2 + \dots + (3 - 3.4)^2$$


$$Error = \sum (r_{ui} - \hat{r}_{ui})^2$$



Singular Value Decomposition



Minimizing the error can lead to overfit



Regularization is mostly used to guard against overfit

Commonly used cost function -

$$Error = \sum (r_{ui} - \hat{r}_{ui})^2 + \lambda (b_u^2 + b_i^2 + ||q_i||^2 + ||p_u||^2)$$

Hyperparameter - λ

Parameters of SVD - p_u , q_i , b_u , and b_i



Non Negative Matrix Factorization

In the context of recommendation engines NMF is very similar to SVD



Only difference

Factors estimated in NMF are non negative

The same cost function is minimized as in the case of SVD



$$\text{Error} = \sum (r_{ui} - \hat{r}_{ui})^2 + \lambda(b_u^2 + b_i^2 + ||q_i||^2 + ||p_u||^2)$$

With constraints that, p_u and q_i are positive




Cost Function

Minimize the cost functions

Stochastic Gradient
Descent

Alternating Least
Squares

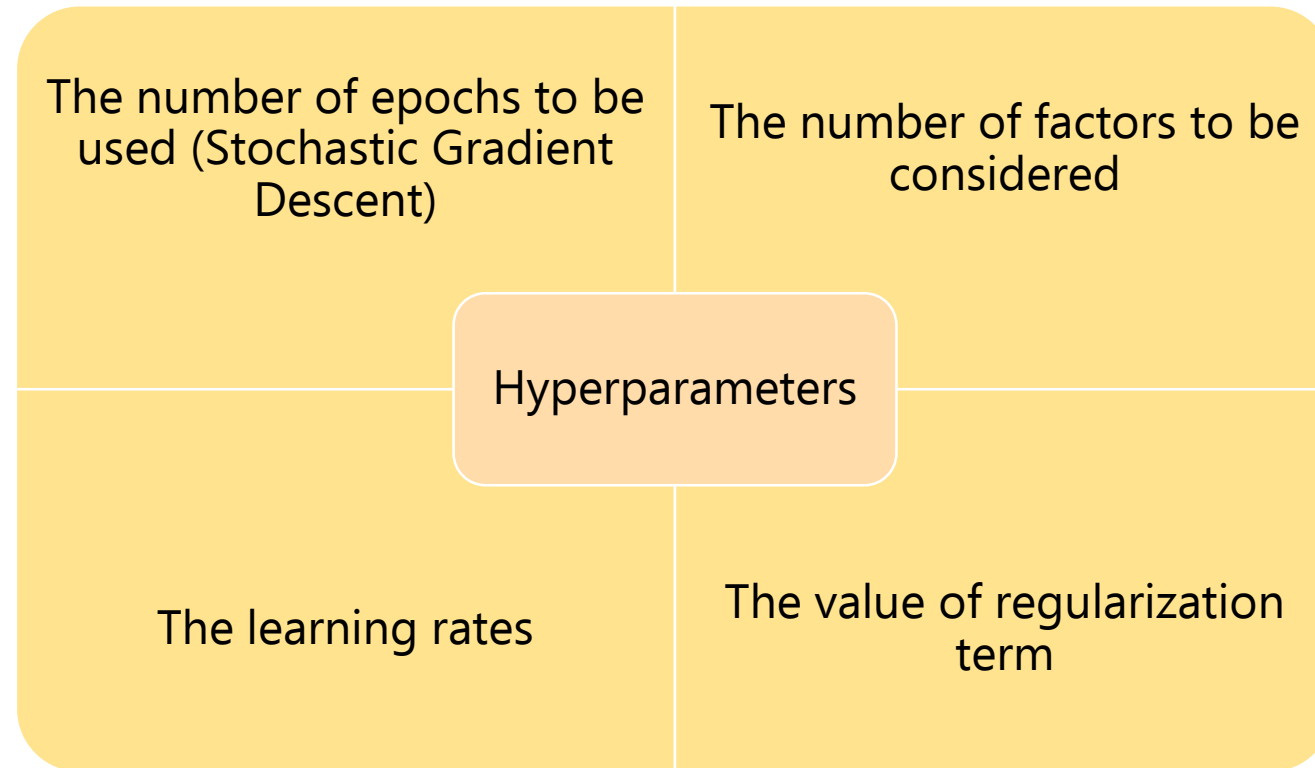


Exact details of implementations can differ depending on the machine learning framework being used

Cost functions discussed here are only indicative of what is normally used by most implementations



Hyperparameters



The optimum values are found out using grid search and cross validation

Recap

- Matrix factorization
- Singular Value Decomposition
- Non Negative Matrix Factorization
- Cost Function
- Hyperparameters

